

CS2109S

AY24/25S2 Midterms

github.com/zeephuru

EXAM INFO

90 minutes, ~ 30 questions

SEAT: MPSH2A, 82

INTELLIGENT AGENTS

PEAS

Performance measure, Environment, Actuators, Sensors

Performance Measure considerations:

- What is best for whom? What are we optimizing? What information is available? Any unintended effects? What are the costs?
- Rational Agent** chooses actions that maximise perf measure.

Task Environment

- Fully** or **Partially** Observable - can observe complete state of environment?
- Deterministic** or **Stochastic** - randomness.
- Strategic** or **Dumb** - prescence of other intelligent agents.
- Episodic** or **Sequential** - episodic: choice of action in atomic episode depends **only** on the episode.
- Static** or **Dynamic** - whether environment changes while agent is deliberating.
- Discrete** or **Continuous** - distinct, clearly defined percepts and actions
- Single** or **Multi-Agent**

Agents

Completely specified by the agent function.

Agent structures

- Simple Reflex Agents** - `if ... then`
- Goal-based Agents** - knows what happens if an action is taken, pick based on goal
- Utility-based Agents** - predicts utility
- Learning Agents** - performance element

Exploration vs. Exploitation

SEARCH

Problem Formulation

- States, Initial State, Goal State/Test: `is_goal(state)`
- Representation Invariant
- Actions: $|actions(state)| \leq b$ (branching factor)
- Transition Model: `new_state = transition(state, action)`
- Action cost function

Criteria

- Completeness
 - Complete $\Leftarrow \forall$ problems, solution exists.
 - Incomplete $\Leftarrow \exists$ problem, solution **does not** exists.
- Optimal
 - Optimal $\Leftarrow \forall$ solution-producing instance, solution is the best.
- Optimal and Incomplete: \exists problem with no solution, but all found solutions are optimal.

Uninformed Search

```
frontier = Frontier()

frontier.append(Node(initial_state))
while frontier is not empty:
    node = frontier.pop()
    if node.state is goal:
        return solution

    for action in actions(node.state):
        next_state = transition(
            node.state, action)
        frontier.add(Node(next_state))

return failure
```

- BFS** - **Queue**, explore layer-by-layer
- UCS** - **Pqueue** (by path cost), creates "tiers" based on costs to reach nodes
- DFS** - **Stack**, go deep then backtrack
- Depth-Limited Search (DLS)** - limit max depth
 - Yes DLS (DFS) is not complete and not optimal
 - "Depth-Limited BFS ensures that we can find the shortest path solution if it exists. The depth limit allows us to terminate in finite time if a solution does not exist."
- Iterative Deepening Search (IDS)** - DLS with incrementing depth limits

Informed Search

- A* Search**
- $$f(n) = g(n) + h(n)$$
- $g(n)$: cost to reach n , $h(n)$: heuristic
 - Frontier: `PQueue(f(n))`
 - Complexities are *exp*
 - Complete if edges costs are positive and b is finite
 - Optimality depends on heuristic

Heuristics

- Admissible**
 - Never over-estimates cost**, it is an optimistic estimate.
 - Theorem: if $h(n)$ admissible, then A* with visited memory is optimal.
 - Relaxed problem**: fewer restrictions. Optimal solution to a **relaxed problem** is an admissible heuristic.
- Consistent**: \forall node N and each successor P ,
 - $h(N) \leq c(N, P) + h(P)$ and $h(G) = 0$
 - Theorem: if $h(n)$ consistent, then A* with visited memory is optimal.
- Dominant**: $\forall n, h_1(n) \geq h_2(n) \implies h_1$ dominates h_2

Local Search

Typically *incomplete* and *suboptimal*

	Perturbatve	<i>Constructive</i>
Search Space	complete solutions	<i>partial solutions</i>
Search step	modify solution	<i>extend solution</i>


Problem Formulation

- States: may not map to actual problem state; represent **potential solutions**
- Initial State, *Goal Test (optional)*

- Successor Function: generate neighbouring states (candidate solutions)
- Evaluation Function**
- Want to **minimize** or **maximize**.

Hill-Climbing

```
curr_state = initial_state
while True:
    best_successor =
        ""a highest-valued successor
        state of curr_state""
    if eval(best_successor) <= eval(curr_state):
        return curr_state
    curr_state = best_successor
```



Adversarial Search

Problem Formulation

- Terminal States**: States of win/lose/draw
- Utility Function**: Value of state from persp of our agent. Typically need to be computed.

Minimax

- Two-player zero-sum game**
- Complete** if tree is finite
 - Optimal** against optimally-playing opponent. Otherwise there will be *faster* strategies.

Alpha-Beta Pruning

```
def minimax(state):
    # find max value reachable from this value
    v = max_value(state, -MAX, MAX)

    # return a next action that has that state
    return action in expand(state) with value v

def max_value(state, a, b):
    if is_terminal(state):
        return utility(state)
    v = -MAX

    # iterate through next states
    for next_state in expand(state):
        v = max(v, min_value(next_state))
        a = max(a, v)
        if v >= b: return v
    return v # max value from current state

def min_value(state, a, b):
    if is_terminal(state):
        return utility(state)
    v = MAX
```

```
for next_state in expand(state):
    v = min(v, max_value(next_state))
    a = min(b, v)
    if v <= a: return v
return v # min value from current state
```

Good move ordering improves pruning effectiveness.

Cutoff

Implying a **cutoff** strategy: halt search halfway and estimate value of midgame states using an **evaluation function**.
Better handles large/infinite game trees.

MACHINE LEARNING

- Unsupervised Learning**: **Unlabeled** data, find patterns/structure
- Supervised Learning**: **Labeled** data, learns mapping
 - Classification**: Predict discrete label or category
 - Regression**: Predict continuous numerical value

Dataset $D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})\}$
True data generating function $f^*(x)$
Hypothesis class H Set of all possible models/functions
 $h: X \rightarrow Y$

Performance Measure

Regression

$$MSE = \frac{1}{N} \sum_{i=1}^N \left(\hat{y}^{(i)} - y^{(i)} \right)^2$$
$$MAE = \frac{1}{N} \sum_{i=1}^N \left| \hat{y}^{(i)} - y^{(i)} \right|$$

Classification

$$Accuracy = \frac{1}{N} \sum_{i=1}^N \mathbf{1}_{\hat{y}^{(i)} = y^{(i)}}$$

		Actual Label	
		Cancer	Benign
Predicted Label	Cancer	2 True Positive	1 False Positive
	Benign	3 False Negative	4 True Negative

- Accuracy**: $TP + TN / total$
- Precision**: $TP / (TP + FP)$
Number of correct out of predicted positives.
- Recall**: $TP / (TP + FN)$
Number predicted out of **actual** positives.
- F1 Score**: $2 \cdot (1/P + 1/R)^{-1}$

Decision Trees

Given n boolean attributes, 2^{2^n} distinct trees.

Entropy:

$$I(P(v_1) \dots P(v_k)) = - \sum_{i=1}^k P(v_i) \log_2 P(v_i)$$

Information Gain:

$$I \text{ before dividing - weighted mean } I \text{ after dividing}$$

$$remainder(A) = \sum_{i=1}^v \frac{p_i + n_i}{p + n} I \left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i} \right)$$

$$IG(A) = I \left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i} \right) - remainder(A)$$

Decision Tree Learning

Greedy, top-down, recursive.

```
def DTL(examples, attributes, default):
    if examples is empty: return default
    if examples have the same classification:
        return classification
    if attributes is empty:
        return mode(examples)
    best = choose_attribute(attributes, examples)
    tree = a new decision tree with root best
    for each value v_i of best:
        examples_i = {rows in examples with best = v_i}
        subtree = DTL(examples_i, attributes - best, mode(examples))
        add a branch to tree with label v_i and subtree subtree
```

Decision trees can have issues of **overfitting** when generalized to new test data.

Occam's Razor: prefer **short/simple** hypotheses over **long/complex**. Less likely to be coincidences.

Pruning

- Min-sample leaf**: Minimum threshold required to be in a leaf node.

If splitting creates a leaf with nodes < threshold, do not split.

Search	Time	Space	Complete	Optimal
BFS	exp	exp	✓	✓
UCS	exp	exp	✓	✓
DFS	exp	m^k	×	×
DLS (DFS)	exp	exp	×	×
IDS	exp	m^k	✓	✓
A*	exp	exp	✓*	<i>depends</i>
Minimax	exp, $O(b^m)$	m^k	✓*	✓*

Time complexity: nodes generated; Space complexity: size of frontier.

- Max-depth**: Note that depth is is of path from **root** to **leaf**
- Data Preprocessing**
- Partition continuous values (binning)
- Deal with missing values: assign common value, drop attribute, drop rows...

LINEAR Regression

Features $x^{(i)} \in \mathbb{R}^d$; target $y^{(i)} \in \mathbb{R}$

Linear model:

$$h_w(x) = w_0(x_0 : 1) + w_1x_1 + \dots + w_dx_d = w^T x$$

- Feature Transformations**
- Feature Engineering**
 - Add a new feature: $z = x^k, z = \log(x), z = e^x$
 - Feature Scaling**
 - Min-max scaling: $z_i = \frac{x_i - \min(x_i)}{\max(x_i) - \min(x_i)}$
 - This one is to [0, 1]
 - Standardization: $z_i = \frac{x_i - \mu_i}{\sigma_i}$

Loss: MSE

$$J_{MSE}(w) = \frac{1}{2N} \sum_{i=1}^N \left(h_w(x^{(i)}) - y^{(i)} \right)^2$$

Normal Equation

- $$w = (X^T X)^{-1} X^T Y$$
- Slow, inverting requires $O(d^3)$
 - $X^T X$ needs to be **invertible**.

Gradient Descent

$$w_j \xleftarrow{\text{update}} w_j - \gamma \frac{\partial J(w_0, w_1 \dots)}{\partial w_j}$$

Learning rate $\gamma ? 0$, a hyperparameter.

Theorem: A **convex** function has a single global minimum.

Theorem: MSE loss function is **convex** for linear and polynomial regression.

To deal with differently-scaled features:

More text.

- Normalize/Standardize (*or other scaling*)
 - This can actually make GD **faster**.
 - Different learning rate γ_i for each weight.
- Variants**
- Introduce **randomness**, may escape local minima
 - Mini-batch**: Consider subset of training data at a time.
 - Stochastic (SGD)**: Select **one** random data point at a time.

LOGISTIC Regression

target $y \in \{0, 1\}$

Sigmoid/Logistic Function

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad \left| \quad \sigma'(x) = \sigma(x)(1 - \sigma(x)) \right.$$
$$h_w(x) = \sigma(w_0x_0 + w_1x_1 + \dots + w_dx_d) = \sigma(w^T x)$$

Then compare against **decision threshold**.

Decision boundary: separates classes in feature space.

Non-linearly separable: use non-linear **feature transformations**, e.g. scale to $[-1, 1]$, then x^2 .

Loss: BCE

True value y , prediction \hat{y} :

$$BCE(y, \hat{y}) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$
$$J_{BCE}(w) = \frac{1}{N} \sum_{i=1}^N BCE \left(y^{(i)}, h_w(x^{(i)}) \right)$$

Theorem: BCE loss function is **convex** for logistic regression.

Multi-class

- target $y \in \{0, 1, 2, \dots, C\}$
- One-vs-one**
- Separate classifier for every pair.
 - Class with most votes selected.
- One-vs-rest**

- Separate classifier for each class against **all other** classes.
- Class with highest probability (confidence score) selected.

Supervised Learning ++

Dataset

- Quality**
- relevance, noise, balance (classification)
- Quantity**: more usually \rightarrow better, but containing **all** is simply **memorization**.

Model Complexity

- Size and expressiveness. e.g. **Polynomial** higher complexity than **linear**.
- Simple model** good enough for simple truth, few data points needed
 - high bias, low variance*
 - Complex model** overfits if **few** data points given
 - low bias, high variance*
 - Bias**: Dependency to fit what it's capable of, **complex model** can fit data well \rightarrow **low bias**.
 - Variance**: How much model changes as **number of data points** changes.

Hyperparameters

Predefined and adjusted manually. By comparison **parameters**, e.g. weights, are learned during training.

- learning rate, feature transformations, batch size/iterations in mini-batch

Hyperparmater Tuning

- Grid Search (exhaustive search): trying all possible combinations
- Random Search: randomly sampling rates and Hyperparameters
- Local Search: e.g. Hill climb

MISC

$$\log_b(a) = \frac{\log_x(a)}{\log_x(b)}$$